

# Linear Regression

# Linear Regression

Consider a (training) dataset with attributes  $X$  and target function  $y$  where  $X$  is an  $n$ -by- $a$  array and  $y$  is an  $n$ -by- $m$  array (usually  $m=1$ )

Let's suppose there are constants  $W$  and  $b$  such that

$$XW + b = y$$

If we append a column containing only 1's to  $X$ , we can simplify to:

$$X_1 W = y$$

From linear algebra, if  $X_1$  is an invertible matrix

$$(X_1)^{-1} X_1 W = (X_1)^{-1} y$$

$$W = (X_1)^{-1} y$$

The we can use  $W$  to predict  $y$  for test examples

# Linear Regression

From linear algebra, if  $X_1$  is an invertible matrix

$$(X_1)^{-1} X_1 W = (X_1)^{-1} y$$

$$W = (X_1)^{-1} y$$

Then we can use  $W$  to predict  $y$  for test examples

However, for most (all?) training sets,  $X_1$  will not be invertible

Nevertheless, there is an approximation called the pseudo-inverse of  $X_1$ , denoted by  $(X_1)^*$

$$(X_1)^* = (X_1^T X_1)^{-1} X_1^T$$

where  $X_1^T$  is the transpose of  $X_1$

then, for we can predict  $y$  in the test set as:

$$y = X_{1\text{test}} W$$

# Linear Regression for Classification

In order to use a regressor (any regressor, not just linear regression) we use a one-hot representation of the output.

Suppose  $y$  is a 1D array of length  $n$  of integers containing the classes in a dataset

As usual, we assume classes are integers in the  $0, \dots, c-1$  range

Then the one-hot representation of  $y$  is an  $n$ -by- $c$  array, where row  $i$  consists of 0 everywhere except at position  $y[i]$ .

For example if  $y = [4, 2, 3, 1, 0, 4]$  and classes are  $0, \dots, 4$ ,  $\text{onehot}(y)$  is given by:

```
[[0., 0., 0., 0., 1.],  
 [0., 0., 1., 0., 0.]  
 [0., 0., 0., 1., 0.],  
 [0., 1., 0., 0., 0.],  
 [1., 0., 0., 0., 0.],  
 [0., 0., 0., 0., 1.]]
```

# Linear Regression for Classification

Thus we train the model to predict not  $y$ , but the one-hot representation of  $y$ :

```
model.fit(X_train, onehot(y_train))
```

When we predict, the model will return the predicted one-hot representation of  $y_{\text{test}}$ .

```
p = model.predict(X_test)
```

$p$  will be a 2D array of size  $X_{\text{train}}.\text{shape}[0]$  by  $\text{onehot}(y_{\text{train}}).\text{shape}[1]$

How do we convert to a single prediction?

We could find the Euclidean distance from each row in  $p[i]$  in  $p$  to the one-hot representation of each of the classes and assign example  $i$  to the class with the most similar representation.

This is equivalent to assigning example  $i$  to the class corresponding to the position of the largest item in  $p[i]$  (that is,  $\text{pred}[i] = \text{argmax}(p[i])$ ).

# Linear Regression for Classification

For example, the following array would result in predictions [4,2,3,1,0,4]:

```
[[0.79115577, 0.63147976, 0.39390119, 0.54309383, 1.07130847],  
 [0.62835492, 0.02026977, 1.81213046, 0.0330209 , 0.38278168],  
 [0.53767076, 0.77982095, 0.91745407, 1.84388367, 0.04652723],  
 [0.85631156, 1.76228783, 0.91840274, 0.15162574, 0.0680549 ],  
 [1.29367029, 0.01124145, 0.63691936, 0.01732445, 0.05142839],  
 [0.16763345, 0.83554417, 0.7718447 , 0.17421716, 1.51564816]])
```

# Examining the models learned by linear regression

```
X = np.load('particles_X.npy')
y = np.load('particles_y.npy')
X_train, X_test, y_train, y_test = split_train_test(X,y)
model = linear_regression()
model.fit(X_train, y_train)
print(model.W)

[[ 0.01351843]
 [-0.04888286]
 [ 0.06925198]
 [-0.16791558]
 [ 0.70958567]
 [-0.72275243]]
```

# Examining the models learned by linear regression

W is a 1x6 vector

The predicted target value of training example x is:

$$p = w[0]*x[0] + \dots + w[4]*x[4] + w[5]$$

The prediction is a linear combination of the attributes

Which attribute is the most important?

```
print(model.W)
```

```
[[ 0.01351843]
```

```
 [-0.04888286]
```

```
 [ 0.06925198]
```

```
 [-0.16791558]
```

```
 [ 0.70958567]    - Coefficient with largest absolute value – Is attribute 4 the most important?
```

```
 [-0.72275243]]   - Constant value added to all predictions
```



# Examining the models learned by linear regression

- Is the attribute with the largest absolute coefficient the most important?
- Not necessarily. Attributes with small values will tend to have larger coefficients and vice versa
- Linear regression is independent to scaling of its features. Thus we can scale the features to make them have unit standard deviation, then fit a linear model
- Then the coefficient with the largest absolute value will normally correspond to the most important attribute (excluding the last element of  $W$ , which is the constant or bias term)

# Beyond Linear Features with Linear Regression

Is it possible to learn a quadratic function such as

$$y = 3x^2 - 5x + 10$$

using linear regression?

# Beyond Linear Features with Linear Regression

For the particle dataset, what if  $y$  depends on the square of one or some of the attributes in  $X$ ?

# Beyond Linear Features with Linear Regression

For the particle dataset, what if  $y$  depends on the square of one or some of the attributes in  $X$ ?

Can we learn a function such as:

$$p = w[0]*x[0] + \dots + w[4]*x[4] + w[5]*x[0]**2 + \dots + w[9]*x[4]**2 + w[10]$$

# Beyond Linear Features with Linear Regression

For the particle dataset, what if  $y$  depends on the product(s) of two of the attributes in  $X$ ?

Can we learn a function such as:

$$p = w[0]*x[0]*x[0] + w[1]*x[0]*x[1] + w[2]*x[0]*x[2] + \dots + w[14]*x[4]*x[4] + w[15]$$

# Beyond Linear Features with Linear Regression

For the particle dataset, what if  $y$  depends on the product(s) of two of the attributes in  $X$ ?

The linear regression algorithm can only learn linear combinations of the attributes

However, we can create new attributes to augment the original dataset with non-linear combinations of the original attributes

# Beyond Linear Features with Linear Regression

We can create new attributes to augment the original dataset with non-linear combinations of the original attributes

What does the following function do?

```
X = np.load('particles_X.npy')
y = np.load('particles_y.npy')
X = np.hstack((X, X**2))
X_train, X_test, y_train, y_test = split_train_test(X, y)
model = linear_regression()
model.fit(X_train, y_train)
pred = model.predict(X_test)
```