


# Convolutional Neural Networks


# Convolutional Neural Networks

- We know it is good to learn a small model.
- From this fully connected model, do we really need all the connections?
- Can some of these be shared?



# Introduction

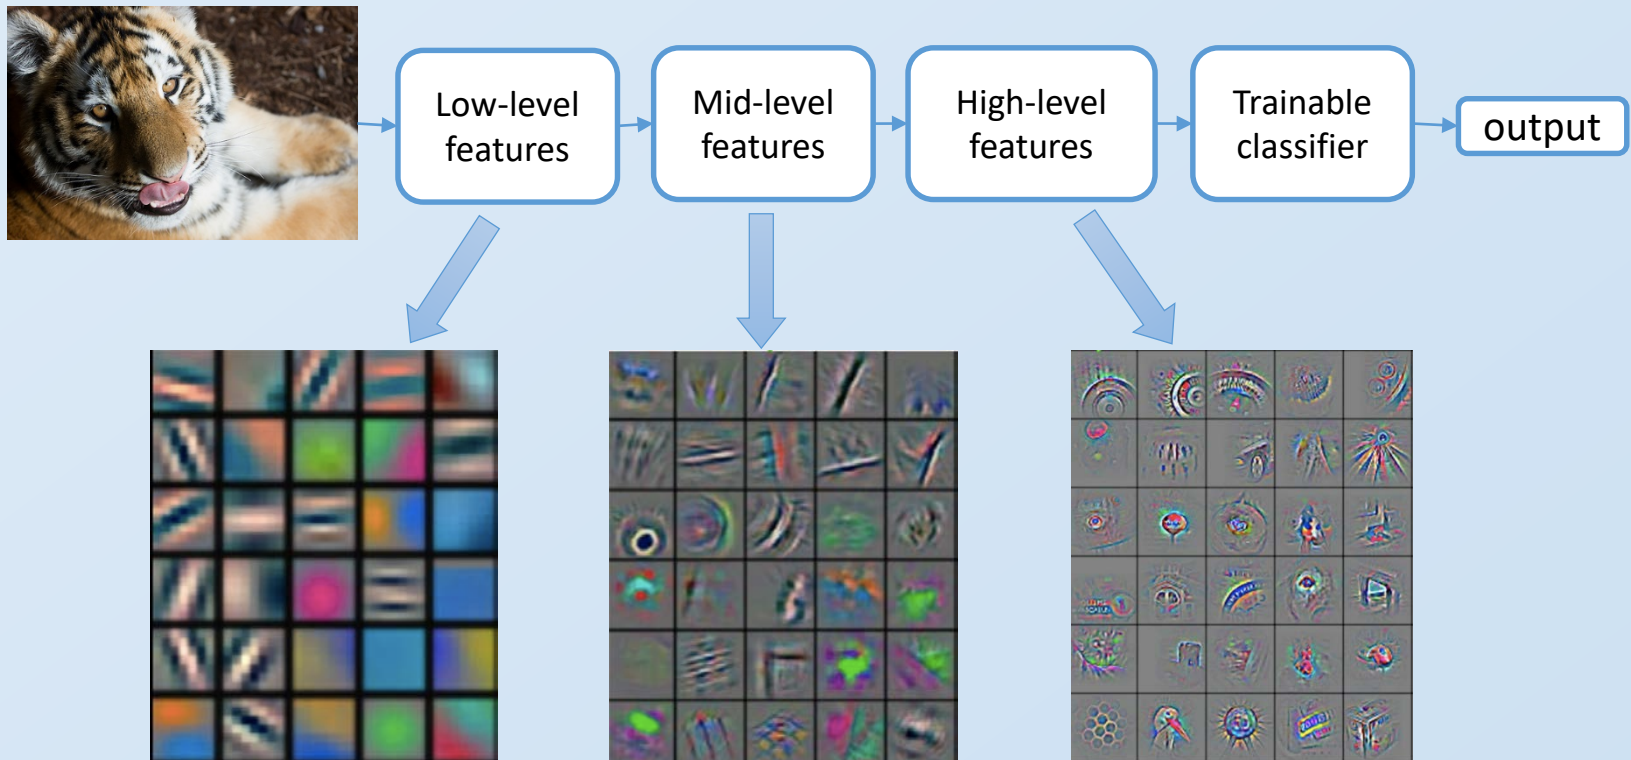
- Traditional pattern recognition models use hand-crafted features and relatively simple trainable classifiers.



- This approach has the following limitations:
  - It is very tedious and costly to develop hand-crafted features
  - The hand-crafted features are usually highly dependent on one application, and cannot be transferred easily to other applications

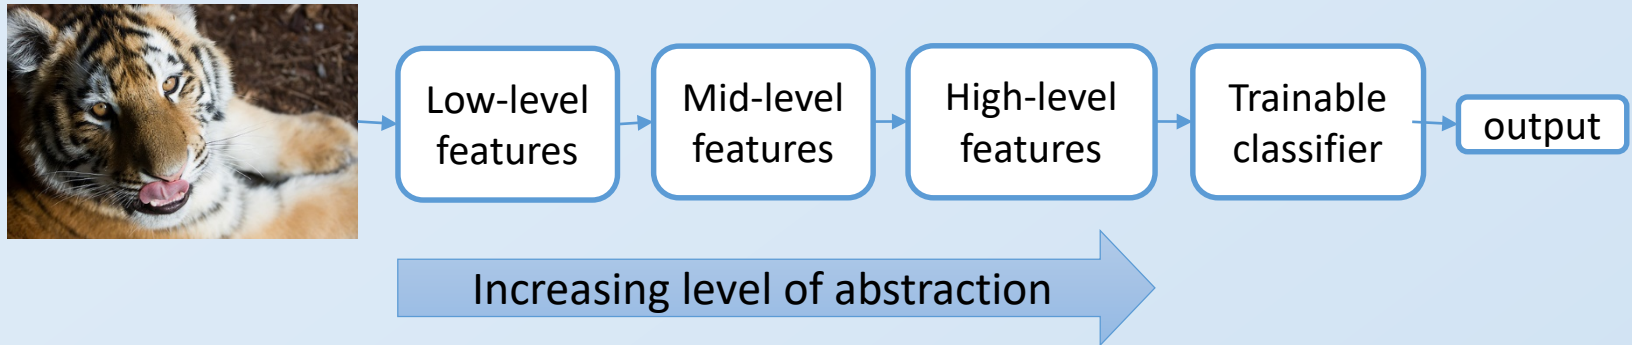
# Deep Learning

- Deep learning (a.k.a. representation learning) seeks to learn rich hierarchical representations (i.e. features) automatically through multiple stages of feature learning.



Feature visualization of convolutional net trained on ImageNet  
(Zeiler and Fergus, 2013)

# Learning Hierarchical Representations




- Hierarchy of representations with increasing level of abstraction. Each stage is a kind of trainable nonlinear feature transform
- Image recognition
  - Pixel  $\rightarrow$  edge  $\rightarrow$  texon  $\rightarrow$  motif  $\rightarrow$  part  $\rightarrow$  object
- Text
  - Character  $\rightarrow$  word  $\rightarrow$  word group  $\rightarrow$  sentence  $\rightarrow$  paragraph  $\rightarrow$  story
- Speech
  - Sound  $\rightarrow$  phoneme  $\rightarrow$  syllable  $\rightarrow$  word  $\rightarrow$  word group  $\rightarrow$  sentence

# Consider learning an image:

- Some patterns are much smaller than the whole image



Can represent a small region with fewer parameters






Same pattern appears in different places:  
They can be compressed!

What about training a lot of such “small” detectors  
and each detector must “move around”.



# A convolutional layer

A CNN is a neural network with some convolutional layers (and some other layers). A convolutional layer has a number of filters that does convolutional operation.





# Convolution

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image

**These are the network parameters to be learned.**

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

-1	1	-1
-1	1	-1
-1	1	-1

Filter 2

⋮ ⋮

Each filter detects a small pattern (3 x 3).

# Convolution

stride=1

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

Dot  
product



3

-1

# Convolution

If stride=2

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

3

-3

# Convolution

stride=1

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

3	-1	-3	-1
-3	1	0	-3
-3	-3	0	1
3	-2	-2	-1

# Convolution

stride=1


1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image

-1	1	-1
-1	1	-1
-1	1	-1

Filter 2

Repeat this for each filter



Two 4 x 4 images  
Forming 2 x 4 x 4 matrix

## Color image




The diagram shows a 6x6 grid of cells. The top layer is a 6x6 grid of white cells. Below it are two more layers of white cells, each shifted one unit to the left and one unit down. The top layer has blue numbers in the corners: 1 in the top-left and top-right cells, and 0 in the bottom-left and bottom-right cells. The middle layer has blue numbers in the corners: 1 in the top-left and top-right cells, and 0 in the bottom-left and bottom-right cells. The bottom layer has blue numbers in the corners: 1 in the top-left and top-right cells, and 0 in the bottom-left and bottom-right cells. The grid is surrounded by a light blue border.

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0






# Convolution v.s. Fully Connected



Fully-  
connected



1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0







6 x 6 image

fewer parameters!







6 x 6 image




Fewer parameters

Even fewer parameters




# The whole CNN

cat dog .....



Can repeat many times



# Max Pooling

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

-1	1	-1
-1	1	-1
-1	1	-1

Filter 2

3	-1	-3	-1
-3	1	0	-3
-3	-3	0	1
3	-2	-2	-1

-1	-1	-1	-1
-1	-1	-2	1
-1	-1	-2	1
-1	0	-4	3

# Why Pooling

- Subsampling pixels will not change the object

bird




Subsampling

bird



We can subsample the pixels to make image smaller



fewer parameters to characterize the image




A CNN compresses a fully connected network in two ways:

- Reducing number of connections
- Shared weights on the edges
- Max pooling further reduces the complexity

# Max Pooling

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0


6 x 6 image




Conv



Max  
Pooling




New image  
but smaller



2 x 2 image


Each filter  
is a channel

# The whole CNN



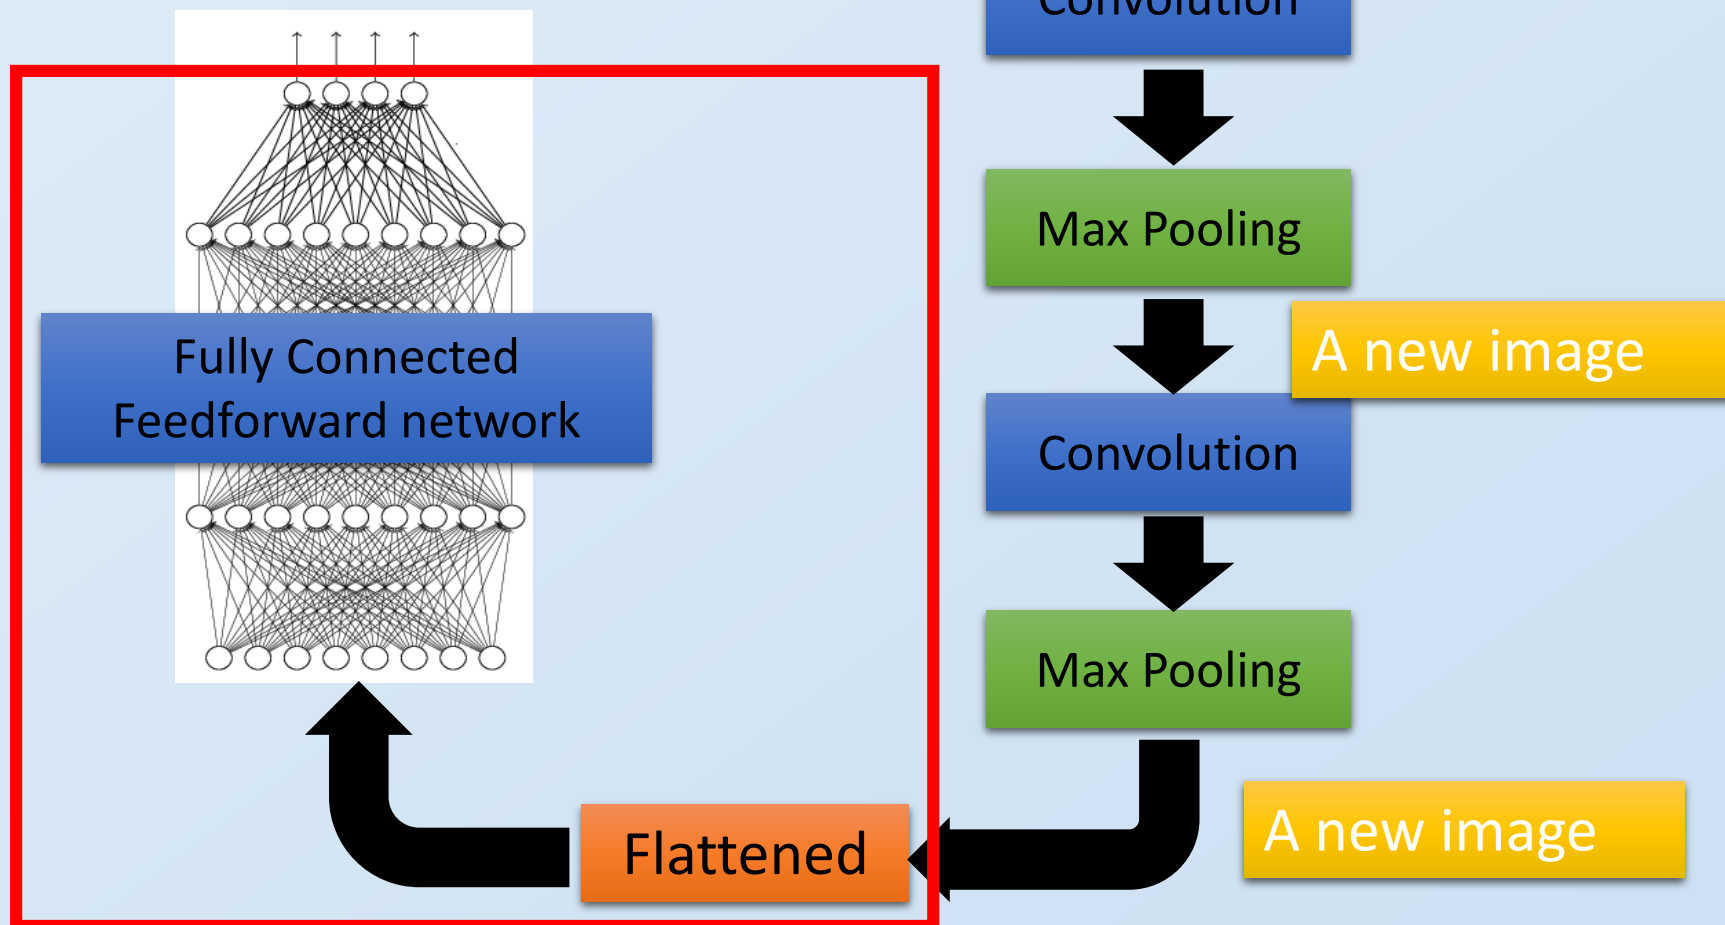
Smaller than the original image

The number of channels is the number of filters




# The whole CNN


cat dog .....



# Flattening



Flattened



Fully Connected  
Feedforward network

# CNN in Keras

Only modified the *network structure* and *input format* (vector -> 3-D tensor)

```
model2.add( Convolution2D( 25, 3, 3,  
                           input_shape=(28, 28, 1)) )
```

1	-1	-1
-1	1	-1
-1	-1	-1

-1	1	-1
-1	1	-1
-1	1	-1




There are  
**25 3x3**  
filters.

Input\_shape = ( 28 , 28 , 1)

28 x 28 pixels

1: black/white, 3: RGB

```
model2.add(MaxPooling2D( (2, 2) ))
```



input

Convolution

Max Pooling

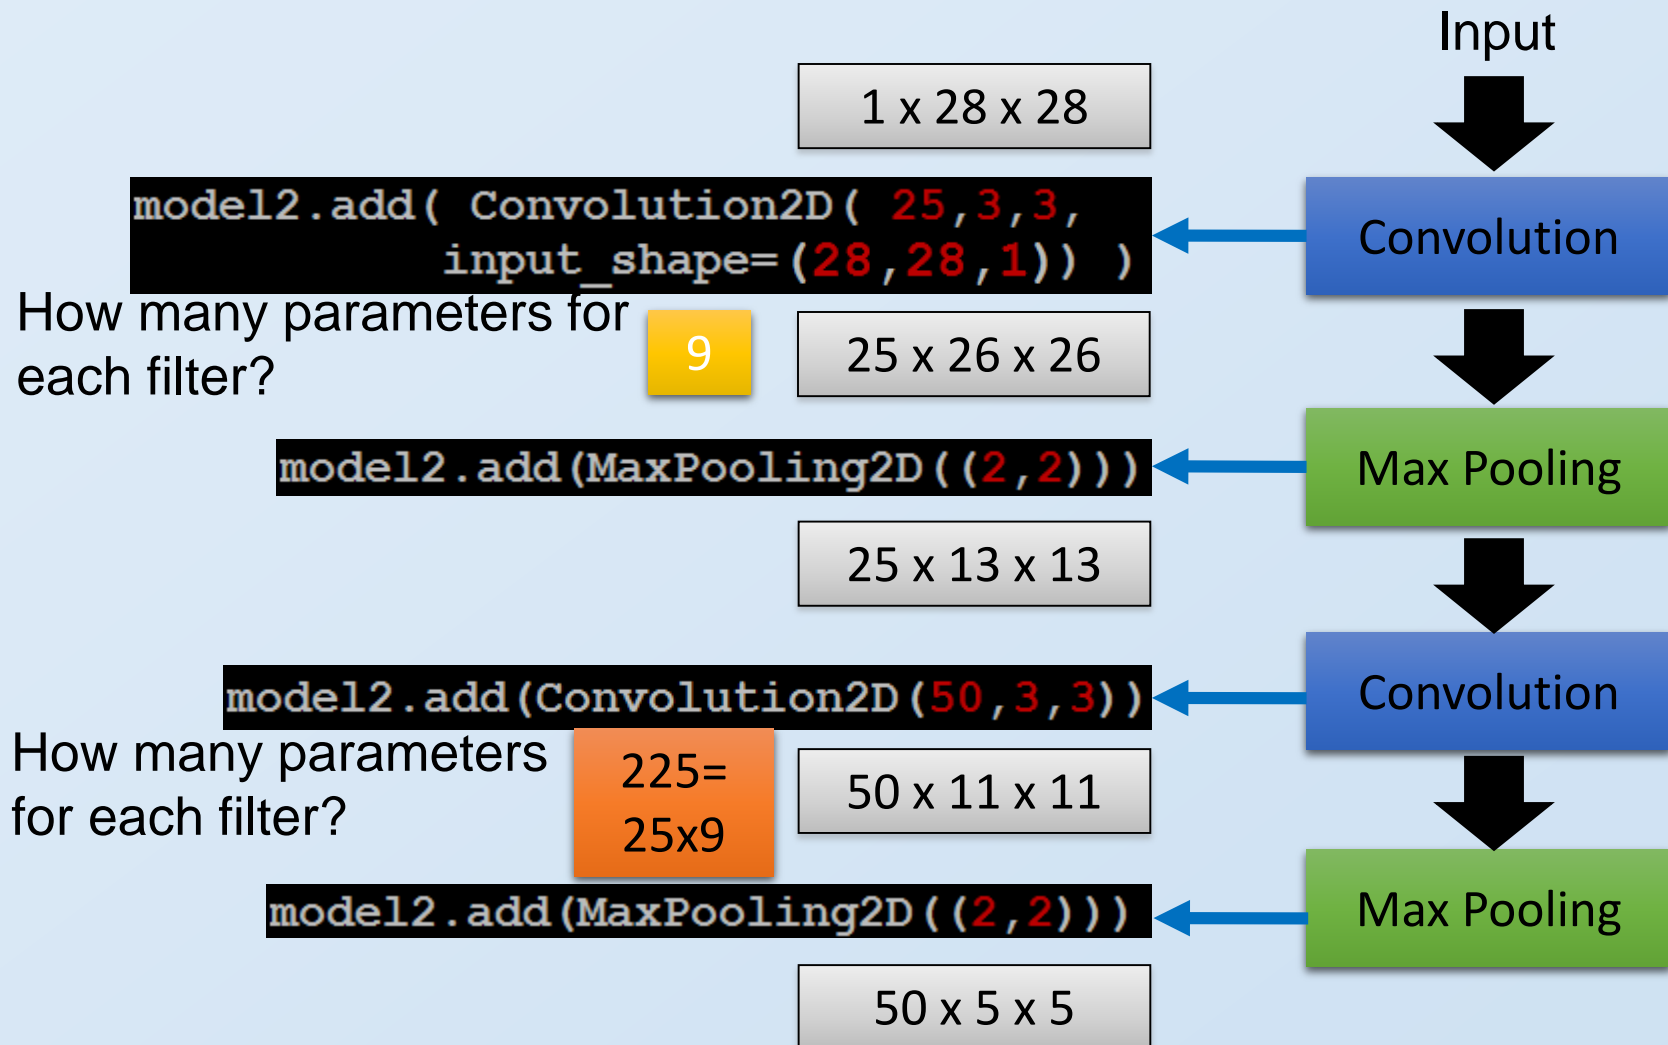
Convolution

Max Pooling



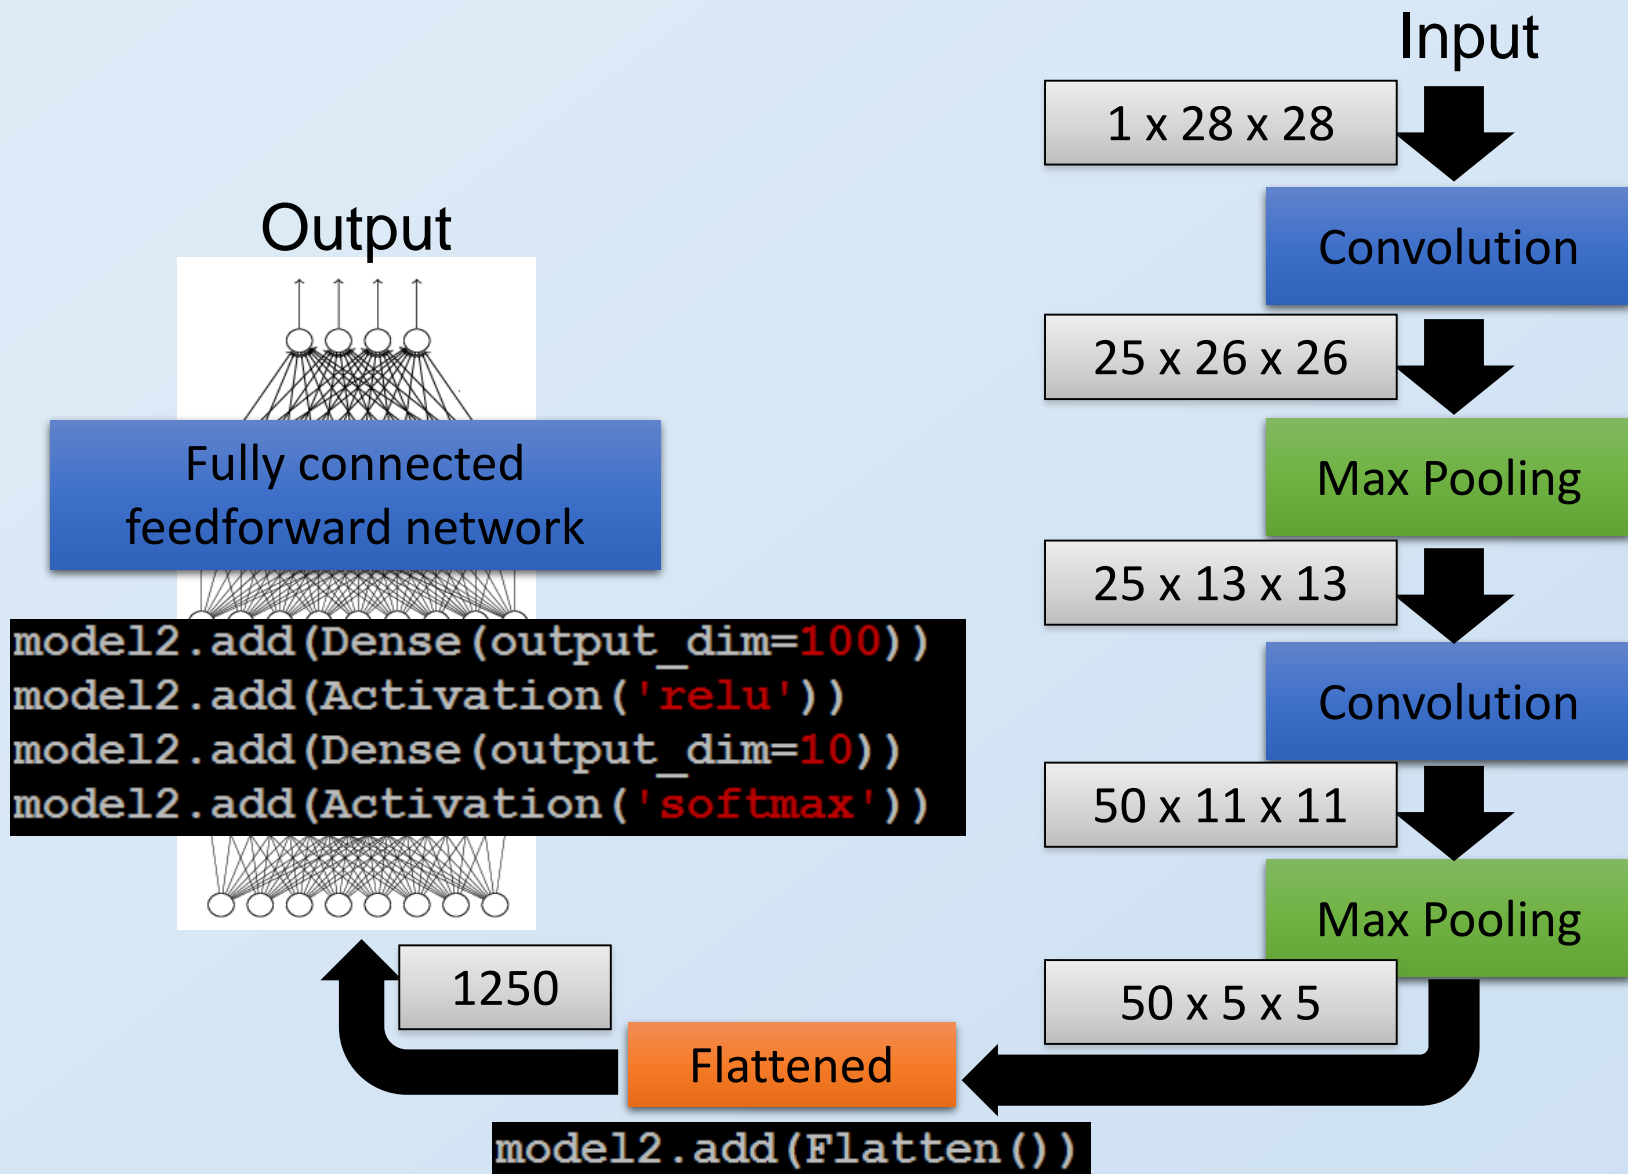
# CNN in Keras

Only modified the *network structure* and *input format* (vector -> 3-D array)

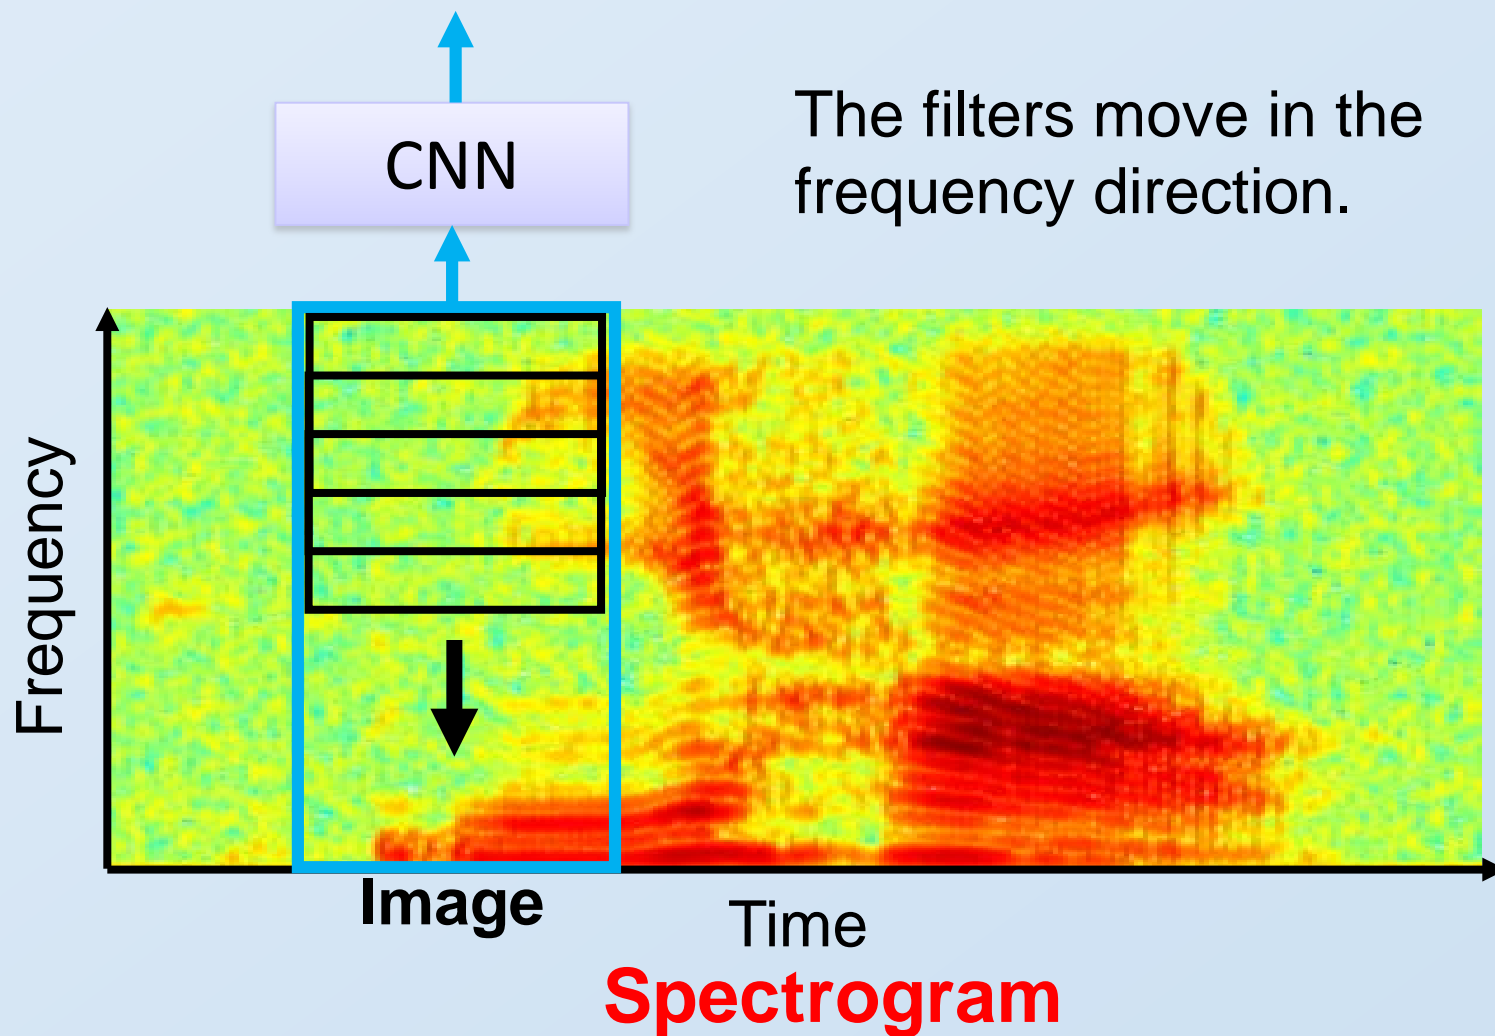


# CNN in Keras

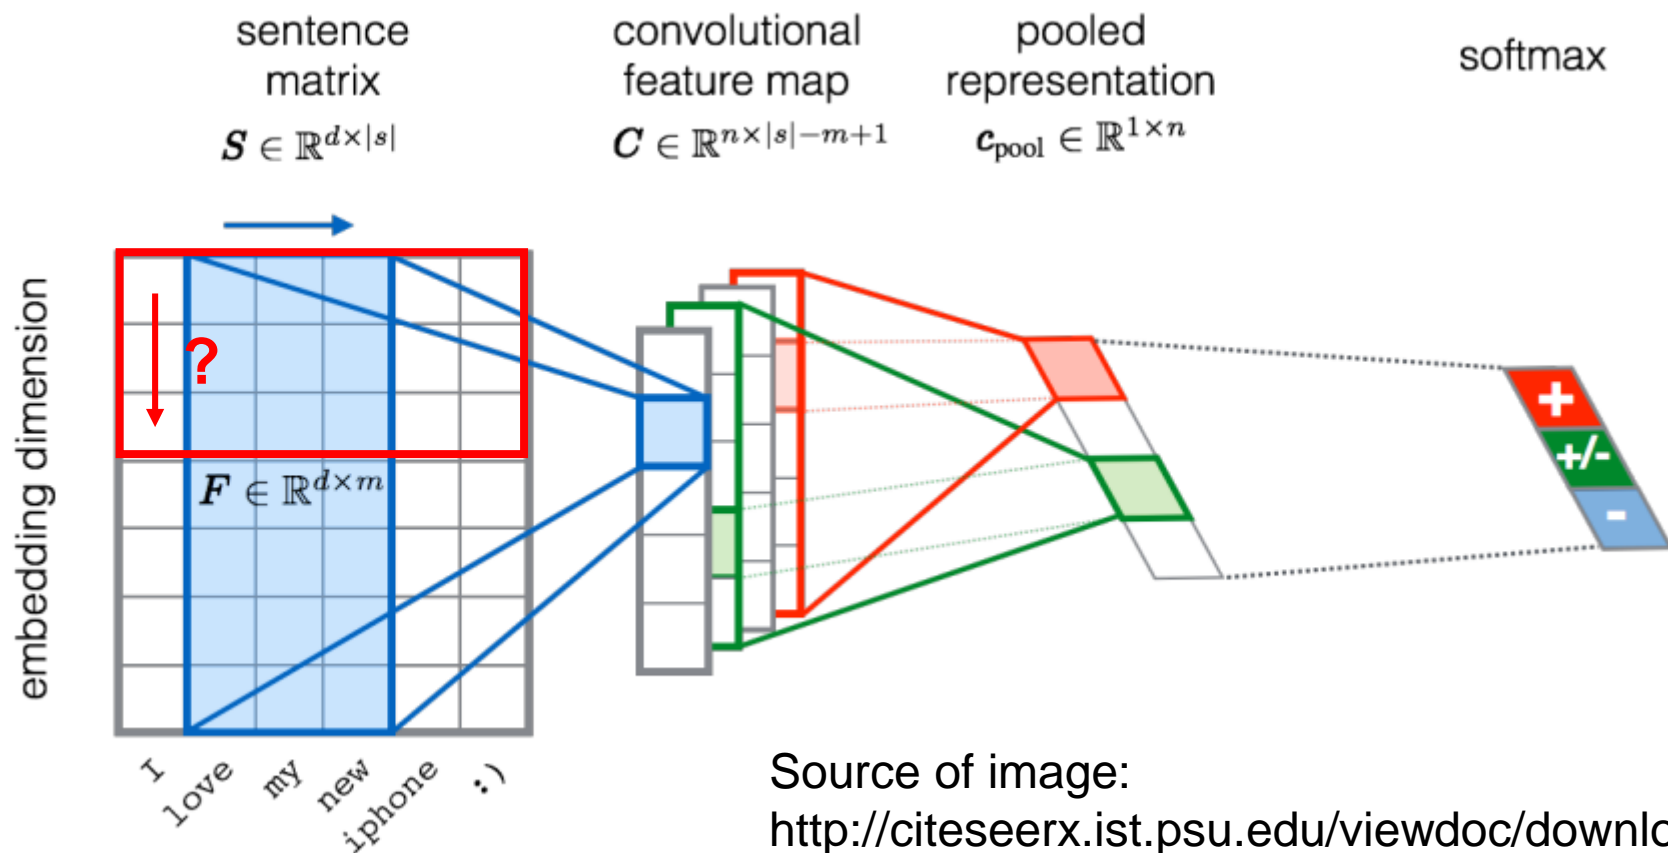
Only modified the *network structure* and *input format (vector -> 3-D array)*



# CNN in speech recognition



# CNN in text classification



Source of image:

<http://citeseerx.ist.psu.edu/viewdoc/download?doi=10.1.1.703.6858&rep=rep1&type=pdf>

# Convolutional Neural Networks

- Questions  
?